


Database Systems

Program in Computer Engineering
Faculty of Engineering

King Mongkut's Institute of Technology Ladkrabang

Text

- Ramez Elmasri and Shamkant B. Navathe.
“**Fundamentals of Database Systems**”
7th Edition., Pearson, 2017



Chapter 7

More SQL: Complex Queries, Views, and Schema
Modification

Outline

- More Complex SQL Retrieval Queries
- Views (Virtual Tables) in SQL
- Schema Modification in SQL

More Complex SQL Retrieval Queries

- Additional features allow users to specify more complex retrievals from database:
 - Nested queries,
 - joined tables, and outer joins (in the FROM clause),
 - aggregate functions, and
 - grouping

Comparisons Involving NULL and Three-Valued Logic

- Meanings of **NULL**
 - Unknown value
 - Unavailable or withheld value
 - Not applicable attribute
- Each individual NULL value considered to be different from every other NULL value
- SQL uses a three-valued logic:
 - **TRUE**, **FALSE**, and **UNKNOWN** (like Maybe)
- **NULL = NULL** comparison **is avoided**

Table 7.1 Logical Connectives in Three-Valued Logic

(a)	AND	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	FALSE	UNKNOWN
	FALSE	FALSE	FALSE	FALSE
	UNKNOWN	UNKNOWN	FALSE	UNKNOWN
(b)	OR	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	TRUE	TRUE
	FALSE	TRUE	FALSE	UNKNOWN
	UNKNOWN	TRUE	UNKNOWN	UNKNOWN
(c)	NOT			
	TRUE	FALSE		
	FALSE	TRUE		
	UNKNOWN	UNKNOWN		

- SQL allows queries that check whether an attribute value is NULL
 - **IS** or **IS NOT NULL**

Query 18. Retrieve the names of all employees who do not have supervisors.

Q18: **SELECT** Fname, Lname
 FROM EMPLOYEE
 WHERE Super_ssn **IS** NULL;

Nested Queries, Tuples, and Set/Multiset Comparisons

- **Nested queries**
 - Complete select-from-where blocks within WHERE clause of another query
 - **Outer query and nested subqueries**
- Comparison **operator** **IN**
 - Compares value v with a set (or multiset) of values V
 - **Evaluates to TRUE if v is one of the elements in V**

- Use tuples of values in comparisons
 - Place them within parentheses

```
SELECT DISTINCT Essn
FROM WORKS_ON
WHERE (Pno, Hours) IN ( SELECT Pno, Hours
                        FROM WORKS_ON
                        WHERE Essn='123456789' );
```

Select the ESSN of all employees who work the same (project, hours) combination on some project that employee 'John Smit' (whose SNN = '123456789') works on.

- Use other comparison operators to compare a single value v
 - = ANY (or = SOME) operator
 - Returns TRUE if the value v is equal to some value in the set V and is hence equivalent to IN
 - Other operators that can be combined with ANY (or SOME):
>, >=, <, <=, and <>
 - ALL: value must exceed all values from nested query

```
SELECT  Lname, Fname
FROM    EMPLOYEE
WHERE   Salary > ALL ( SELECT  Salary
                       FROM    EMPLOYEE
                       WHERE   Dno=5 );
```

Return the name of employees whose salary is greater than the salary of all the employees in department 5.

- Avoid potential errors and ambiguities
 - Create tuple variables (aliases) for all tables referenced in SQL query

Query 16. Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

```

Q16:  SELECT    E.Fname, E.Lname
      FROM      EMPLOYEE AS E
      WHERE     E.Ssn IN ( SELECT    Essn
                          FROM      DEPENDENT AS D
                          WHERE     E.Fname=D.Dependent_name
                          AND E.Sex=D.Sex );
  
```

Correlated Nested Queries

- **Queries that are nested using the = or IN comparison operator** can be collapsed into one single block: E.g., Q16 can be written as:

Q16A:

```

SELECT  E.Fname, E.Lname
FROM    EMPLOYEE AS E , DEPENDENT AS D
WHERE   E.Ssn = D.Essn
        AND    E.Sex = D.Sex
        AND    E.Fname = D.Dependent_name;

```

- **Correlated** nested query
 - Evaluated once for each tuple in the outer query

The EXISTS and UNIQUE Functions in SQL for correlating queries

- **EXISTS** function
 - Check whether the result of a correlated nested query is empty or not. They are Boolean functions that return a TRUE or FALSE result.
- **EXISTS** and **NOT EXISTS**
 - Typically used in conjunction with a correlated nested query
- SQL function **UNIQUE (Q)**
 - Returns TRUE if there are no duplicate tuples in the result of query Q

Q16B: Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

```

SELECT  Fname, Lname
FROM    Employee AS E
WHERE    EXISTS  ( SELECT *
                     FROM    DEPENDENT AS D
                     WHERE    E.Ssn = D.Essn AND
                               E.Sex = D.Sex  AND
                               E.Firstname = D.Dependent_name );
  
```


Q6: Retrieve the names of employees who have no dependents.

```

SELECT Fname, Lname
FROM Employee AS E
WHERE NOT EXISTS ( SELECT *
                     FROM DEPENDENT
                     WHERE Ssn = Essn );
  
```

Q7: List the names of managers who have at least one dependent.

```

SELECT Fname, Lname
FROM Employee
WHERE EXISTS (
    SELECT *
    FROM DEPENDENT
    WHERE Ssn = Essn )
    AND
    EXISTS (
    SELECT *
    FROM Department
    WHERE Ssn = Mgr_Ssn )

```

Select all DEPENDENT tuples related to an EMPLOYEE

Select all DEPARTMENT tuples managed by the EMPLOYEE

Explicit Sets and Renaming of Attributes in SQL

- Can use explicit set of values in WHERE clause

Q17: **SELECT DISTINCT** Essn
 FROM WORKS_ON
 WHERE Pno IN (1, 2, 3);

- Use qualifier AS followed by desired new name
 - Rename any attribute that appears in the result of a query

Q8A: **SELECT** E.Lname **AS** Employee_name, S.Lname **AS** Supervisor_name
 FROM EMPLOYEE **AS** E, EMPLOYEE **AS** S
 WHERE E.Super_ssn=S.Ssn;

Specifying Joined Tables in the FROM Clause of SQL

- **Joined table**

- Permits users to specify a table resulting from a join operation in the FROM clause of a query

- The FROM clause in Q1A

- Contains a single joined table.

JOIN may also be called INNER JOIN

```
Q1A:  SELECT  Fname, Lname, Address
      FROM    (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)
      WHERE   Dname='Research';
```

Different Types of JOINed Tables in SQL

- Specify different types of join
 - **NATURAL JOIN**
 - Various types of **OUTER JOIN** (**LEFT**, **RIGHT**, **FULL**)
- **NATURAL JOIN** on two relations R and S
 - No join condition specified
 - Is equivalent to an implicit EQUIJOIN condition for each pair of attributes with same name from R and S

NATURAL JOIN

- Rename attributes of one relation so it can be joined with another using NATURAL JOIN:

```

Q1B:      SELECT   Fname, Lname, Address
FROM      ( EMPLOYEE NATURAL JOIN
              ( DEPARTMENT AS
                DEPT (Dname, Dno, Mssn, Msdate) ) )
WHERE     Dname='Research';
  
```

The above works with EMPLOYEE.Dno = DEPT.Dno
as an implicit join condition

INNER and OUTER Joins

- **INNER JOIN** (**versus** OUTER JOIN)
 - Default type of join in a joined table
 - Tuple is included in the result only if a matching tuple exists in the other relation
- **LEFT OUTER JOIN**
 - Every tuple in left table must appear in result
 - If no matching tuple
 - Padded with NULL values for attributes of right table
- **RIGHT OUTER JOIN**
 - Every tuple in right table must appear in result
 - If no matching tuple
 - Padded with NULL values for attributes of left table

Aggregate Functions in SQL

- Used to summarize information from multiple tuples into a single-tuple summary
- Built-in aggregate functions
 - **COUNT**, **SUM**, **MAX**, **MIN**, and **AVG**
- **Grouping**
 - Create subgroups of tuples before summarizing
- To select entire groups, **HAVING** clause is used
- Aggregate functions can be used in the **SELECT** clause or in a **HAVING** clause

Renaming Results of Aggregation

- Following query returns a single row of computed values from EMPLOYEE table:

Q19: **SELECT** **SUM** (Salary), **MAX** (Salary), **MIN** (Salary), **AVG** (Salary)
 FROM EMPLOYEE;

- The result can be presented with new names:

Q19A: **SELECT** **SUM** (Salary) **AS** Total_Sal, **MAX** (Salary) **AS** Highest_Sal,
 MIN (Salary) **AS** Lowest_Sal, **AVG** (Salary) **AS** Average_Sal
 FROM EMPLOYEE;

- NULL values are discarded when aggregate functions are applied to a particular column

Query 20. Find the sum of the salaries of all employees of the 'Research' department, as well as the maximum salary, the minimum salary, and the average salary in this department.

```
Q20:  SELECT    SUM (Salary), MAX (Salary), MIN (Salary), AVG (Salary)
      FROM      (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)
      WHERE     Dname='Research';
```

Queries 21 and 22. Retrieve the total number of employees in the company (Q21) and the number of employees in the 'Research' department (Q22).

```
Q21:  SELECT    COUNT (*)
      FROM      EMPLOYEE;
```

```
Q22:  SELECT    COUNT (*)
      FROM      EMPLOYEE, DEPARTMENT
      WHERE     DNO=DNUMBER AND DNAME='Research';
```

Aggregate Functions on Booleans

- SOME and ALL may be applied as functions on Boolean Values.
 - **SOME** returns true if **at least one element** in the collection is TRUE (similar to OR)
 - **ALL** returns true if **all of the elements** in the collection are TRUE (similar to AND)

Grouping: The GROUP BY Clause

- **Partition** relation into subsets of tuples
 - Based on **grouping attribute(s)**
 - Apply function to each such group independently
- **GROUP BY clause**
 - Specifies grouping attributes
- **COUNT (*)** counts the number of rows in the group

Examples of GROUP BY

- The grouping attribute must appear in the SELECT clause:

Q24: **SELECT** Dno, **COUNT** (*), **AVG** (Salary)
 FROM EMPLOYEE
 GROUP BY Dno;

- If the grouping attribute has NULL as a possible value, then a separate group is created for the null value (e.g., null Dno in the above query)
- GROUP BY may be applied to the result of a JOIN:

Q25: **SELECT** Pnumber, Pname, **COUNT** (*)
 FROM PROJECT, WORKS_ON
 WHERE Pnumber=Pno
 GROUP BY Pnumber, Pname;

- **HAVING clause**
 - Provides a condition to select or reject an entire group:
- **Query 26.** For each project *on which more than two employees work*, retrieve the project number, the project name, and the number of employees who work on the project.

Q26:

SELECT	Pnumber, Pname, COUNT (*)
FROM	PROJECT, WORKS_ON
WHERE	Pnumber=Pno
GROUP BY	Pnumber, Pname
HAVING	COUNT (*) > 2;

Pname	Pnumber	...	Essn	Pno	Hours
ProductX	1		123456789	1	32.5
ProductX	1		453453453	1	20.0
ProductY	2		123456789	2	7.5
ProductY	2		453453453	2	20.0
ProductY	2		333445555	2	10.0
ProductZ	3		666884444	3	40.0
ProductZ	3		333445555	3	10.0
Computerization	10		333445555	10	10.0
Computerization	10		999887777	10	10.0
Computerization	10		987987987	10	35.0
Reorganization	20		333445555	20	10.0
Reorganization	20		987654321	20	15.0
Reorganization	20		888665555	20	NULL
Newbenefits	30		987987987	30	5.0
Newbenefits	30		987654321	30	20.0
Newbenefits	30		999887777	30	30.0

These groups are not selected by the HAVING condition of Q26.

After applying the WHERE clause but before applying HAVING

Pname	<u>Pnumber</u>	...	<u>Essn</u>	<u>Pno</u>	Hours	
ProductY	2		123456789	2	7.5	→
ProductY	2		453453453	2	20.0	
ProductY	2		333445555	2	10.0	
Computerization	10		333445555	10	10.0	→
Computerization	10	...	999887777	10	10.0	
Computerization	10		987987987	10	35.0	
Reorganization	20		333445555	20	10.0	→
Reorganization	20		987654321	20	15.0	
Reorganization	20		888665555	20	NULL	
Newbenefits	30		987987987	30	5.0	→
Newbenefits	30		987654321	30	20.0	
Newbenefits	30		999887777	30	30.0	

Pname	Count (*)
ProductY	3
Computerization	3
Reorganization	3
Newbenefits	3

Result of Q26
(Pnumber not shown)

After applying the HAVING clause condition

Combining the WHERE and the HAVING Clause

- Consider the query: we want to count the *total* number of employees whose salaries exceed \$40,000 in each department, but only for departments where more than five employees work.

- INCORRECT QUERY:**

```

SELECT      Dno, COUNT (*)
FROM        EMPLOYEE
WHERE       Salary > 40000
GROUP BY    Dno
HAVING      COUNT (*) > 5;
  
```

Correct Specification of the Query:

- Note: the WHERE clause applies tuple by tuple whereas HAVING applies to entire group of tuples

Query 28. For each department that has more than five employees, retrieve the department number and the number of its employees who are making more than \$40,000.

```
Q28:  SELECT  Dnumber, COUNT (*)
      FROM    DEPARTMENT, EMPLOYEE
      WHERE   Dnumber=Dno AND Salary>40000 AND
             ( SELECT      Dno
               FROM        EMPLOYEE
               GROUP BY Dno
               HAVING      COUNT (*) > 5)
```

EXPANDED Block Structure of SQL Queries

```

SELECT <attribute and function list>
FROM <table list>
[ WHERE <condition> ]
[ GROUP BY <grouping attribute(s)> ]
[ HAVING <group condition> ]
[ ORDER BY <attribute list> ];
  
```

Views (Virtual Tables) in SQL

- Concept of a view in SQL
 - Single table derived from other tables called the **defining tables**
 - Considered to be a virtual table that is not necessarily populated

Specification of Views in SQL

- **CREATE VIEW** command

- Give table name, list of attribute names, and a query to specify the contents of the view
- In V1, attributes retain the names from base tables. In V2, attributes are assigned names

V1:	CREATE VIEW	WORKS_ON1
	AS SELECT	Fname, Lname, Pname, Hours
	FROM	EMPLOYEE, PROJECT, WORKS_ON
	WHERE	Ssn=Essn AND Pno=Pnumber;
V2:	CREATE VIEW	DEPT_INFO(Dept_name, No_of_emps, Total_sal)
	AS SELECT	Dname, COUNT (*), SUM (Salary)
	FROM	DEPARTMENT, EMPLOYEE
	WHERE	Dnumber=Dno
	GROUP BY	Dname;

- Once a View is defined, SQL queries can use the View relation in the FROM clause
- View is always up-to-date
 - Responsibility of the DBMS and not the user
- **DROP VIEW** command
 - Dispose of a view

View Implementation, View Update, and Inline Views

- Complex problem of efficiently implementing a view for querying
- **Strategy1: Query modification** approach
 - Compute the view as and when needed. Do not store permanently
 - Modify view query into a query on underlying base tables
 - **Disadvantage:** inefficient for views defined via complex queries that are time-consuming to execute

- **Strategy 2: View materialization**

- Physically create a temporary view table when the view is first queried
- Keep that table on the assumption that other queries on the view will follow
- Requires efficient strategy for automatically updating the view table when the base tables are updated

- **Incremental update strategy for materialized views**

- DBMS determines what new tuples must be inserted, deleted, or modified in a materialized view table

- Multiple ways to handle materialization:
 - **immediate update** strategy updates a view as soon as the base tables are changed
 - **lazy update** strategy updates the view when needed by a view query
 - **periodic update** strategy updates the view periodically (in the latter strategy, a view query may get a result that is not up-to-date). This is commonly used in Banks, Retail store operations, etc.

View Update

- Update on a view defined on a single table without any aggregate functions
 - Can be mapped to an update on underlying base table?
 - possible if the primary key is preserved in the view
- Update not permitted on aggregate views. E.g.,

UV2: UPDATE	DEPT_INFO
SET	Total_sal=100000
WHERE	Dname='Research';

cannot be processed because **Total_sal is a computed value** in the view definition

- View involving joins
 - Often not possible for DBMS to determine which of the updates is intended
- Clause **WITH CHECK OPTION**
 - Must be added at the end of the view definition if a view is to be updated to make sure that tuples being updated stay in the view

Views as authorization mechanism

- SQL query authorization statements (GRANT and REVOKE) are described in detail in Chapter 30
- Views can be used to hide certain attributes or tuples from unauthorized users
- E.g., For a user who is only allowed to see employee information for those who work for department 5, he may only access the view **DEPT5EMP**:

```
CREATE VIEW      DEPT5EMP AS
SELECT          *
FROM            EMPLOYEE
WHERE           Dno = 5;
```

Schema Change Statements in SQL

- **Schema evolution commands**
 - DBA may want to change the schema while the database is operational
 - Does not require recompilation of the database schema

The DROP Command

- **DROP** command
 - Used to drop named schema elements, such as tables, domains, or constraint
- Drop behavior options:
 - CASCADE and RESTRICT
- Example:
 - `DROP SCHEMA COMPANY CASCADE;`
 - This removes the schema and all its elements including tables, views, constraints, etc.

The ALTER table command

- **Alter table actions** include:
 - Adding or dropping a column (attribute)
 - Changing a column definition
 - Adding or dropping table constraints
- **Example:**
 - `ALTER TABLE COMPANY.EMPLOYEE ADD COLUMN Job
VARCHAR(12) ;`

Adding and Dropping Constraints

- Change constraints specified on a table
 - Add or drop a named constraint

```
ALTER TABLE COMPANY.EMPLOYEE
DROP CONSTRAINT EMPSUPERFK CASCADE;
```


Dropping Columns, Default Values

- To drop a column
 - Choose either CASCADE or RESTRICT
 - CASCADE would drop the column from views etc.
RESTRICT is possible if no views refer to it.

ALTER TABLE COMPANY.EMPLOYEE DROP COLUMN Address CASCADE;

- Default values can be dropped and altered :

ALTER TABLE COMPANY.DEPARTMENT ALTER COLUMN Mgr_ssn DROP DEFAULT;
ALTER TABLE COMPANY.DEPARTMENT ALTER COLUMN Mgr_ssn SET DEFAULT
'333445555';

Table 7.2 Summary of SQL Syntax

Table 7.2 Summary of SQL Syntax

<p>CREATE TABLE <table name> (<column name> <column type> [<attribute constraint>] { , <column name> <column type> [<attribute constraint>] } [<table constraint> { , <table constraint> }])</p>
<p>DROP TABLE <table name></p>
<p>ALTER TABLE <table name> ADD <column name> <column type></p>
<p>SELECT [DISTINCT] <attribute list> FROM (<table name> { <alias> } <joined table>) { , (<table name> { <alias> } <joined table>) } [WHERE <condition>] [GROUP BY <grouping attributes> [HAVING <group selection condition>]] [ORDER BY <column name> [<order>] { , <column name> [<order>] }]</p>
<p><attribute list> ::= (* (<column name> <function> (([DISTINCT] <column name> *))) { , (<column name> <function> (([DISTINCT] <column name> *))) }))</p>
<p><grouping attributes> ::= <column name> { , <column name> }</p>
<p><order> ::= (ASC DESC)</p>
<p>INSERT INTO <table name> [(<column name> { , <column name> })] (VALUES (<constant value> , { <constant value> }) { , (<constant value> { , <constant value> }) } <select statement>)</p>

Table 7.2 Summary of SQL Syntax

DELETE FROM <table name>

[WHERE <selection condition>]

UPDATE <table name>

SET <column name> = <value expression> { , <column name> = <value expression> }

[WHERE <selection condition>]

CREATE [UNIQUE] INDEX <index name>

ON <table name> (<column name> [<order>] { , <column name> [<order>] })

[CLUSTER]

DROP INDEX <index name>

CREATE VIEW <view name> [(<column name> { , <column name> })]

AS <select statement>

DROP VIEW <view name>

NOTE: The commands for creating and dropping indexes are not part of standard SQL.

Summary

- Complex SQL:
 - Nested queries, joined tables (in the FROM clause), outer joins, aggregate functions, grouping
- **CREATE VIEW** statement and materialization strategies
- Schema Modification for the DBAs using **ALTER TABLE** , **ADD** and **DROP COLUMN** , **ALTER CONSTRAINT** etc.

